Counterfactual Defect Localization of Reward Function Vulnerabilities in Autonomous Driving Decision Modules: A Causal Inference Approach

Wenjia Niu School of Cyberspace Science and Technology
Beijing Jiaotong University, Beijing, China 100044
Email: Niuwj@bjtu.edu.cn

***Abstract*—Reward function defects constitute a critical class of security vulnerabilities in autonomous driving (AD) decision modules: an omitted or mis-specified feature can induce unsafe behaviors that evade traditional software testing, creating latent attack surfaces exploitable by adversarial perturbations of the operating environment. Despite their severity, localizing such defects remains fundamentally challenging because (i) behavioral logs capture only symptoms with no causal attribution to specific reward features, (ii) existing white-box methods require access to proprietary reward source code, and (iii) purely data-driven approaches suffer from high false-positive rates due to confounding. In this paper, we formalize reward defect localization as a *structured causal inference* problem and propose CFD-LOC, a provably convergent black-box framework that treats the AD system as an opaque function mapping scenarios to trajectories. Our key technical innovation is the combination of three formally grounded components: a *counterfactual scenario mutation operator* with guaranteed feature isolation, a *driving invariant library* formalized in linear temporal logic, and a *bayesian hypothesis test* for defect diagnosis with calibrated confidence. We prove that under mild assumptions on scenario coverage and invariant completeness, CFD-LOC achieves exponential convergence of the localization error. Extensive experiments on the Baidu Apollo platform with 8 implanted defect types across 40 test scenarios demonstrate that CFD-LOC achieves 92.5% localization accuracy—substantially outperforming statistical baselines (32.5%), metamorphic testing (45.0%), and LLM-only approaches (55.0%), while requiring only 14.8 min per diagnosis. Ablation studies confirm that each component contributes statistically significantly ($p < 0.001$). To our knowledge, this is the first work to systematically treat reward function security vulnerabilities through the lens of controlled causal experimentation.**

***Index Terms*—Autonomous driving, reward function vulnerabilities, counterfactual inference, invariant testing, black-box defect localization, causal experimentation.**

I. INTRODUCTION

A. Security Context and Motivation

Modern autonomous driving (AD) systems are safety-critical cyber-physical systems whose high-level decision-making is governed by reward or cost functions. These functions encode domain knowledge—ranging from lane-keeping preferences to priorities among safety, comfort, and efficiency—into a numerical optimization objective. However, the complexity of real-world driving scenarios makes reward function design fundamentally error-prone. A designer may inadvertently *omit* a critical feature (e.g., relative velocity), *miscalibrate* competing objectives (e.g., comfort vs. safety), or specify incentives that produce unintended behaviors ("reward hacking" [1]).

From a *security and forensics* perspective, reward function defects represent a distinct class of vulnerabilities:

- They are **latent**: a missing feature may only manifest under specific scenario conditions, remaining dormant during standard testing.
- They are **exploitable**: adversarial perturbations of the environment can trigger reward-induced unsafe behaviors without modifying the system [**?**].
- They are **difficult to attribute**: behavior logs exhibit symptoms (e.g., a collision at a roundabout) but offer no direct evidence of *which* reward feature caused the failure.

B. The Causal Attribution Problem

We formalize the AD system as an opaque function $\mathcal{F} : \mathcal{S} \to \mathcal{T}$, where $\mathcal{S}$ denotes the scenario space (road geometry, vehicle states, pedestrians, environmental conditions) and $\mathcal{T}$ denotes the trajectory space (sequences of ego-vehicle poses and controls). The internal reward function R, policy π, and all intermediate representations are inaccessible—a strict black-box constraint that prohibits code analysis or gradient-based methods.

The black-box defect localization problem is fundamentally a *causal inference* problem: given observed abnormal behavior under a scenario s, how can we determine which feature omission in R causes the violation? Existing approaches fall into three categories with significant limitations:

1) **White-box analysis** [2] requires reward source code access, which is unavailable for proprietary modules.
2) **Metamorphic testing** [3] detects inconsistencies but lacks causal attribution capability, achieving only anomaly-level detection rather than feature-level localization.
3) **LLM-based inference** suffers from high false-positive rates (35.6% in our experiments) due to hallucination in safety-critical judgments.

C. Our Approach: CFD-LOC

We propose CFD-LOC, a *counterfactual causal experimentation* framework that localizes reward feature omissions by systematically constructing minimally-different scenario pairs and observing whether the resulting behavioral change is consistent with the hypothesized missing feature. The framework consists of four formally grounded stages:

1) **Invariant Library ($\mathcal{L}$):** 10 domain-recognized safety and comfort rules formalized in linear temporal logic as objective test oracles.
2) **Anomaly Detection and Hypothesis Generation:** Behavioral deviations from $\mathcal{L}$ trigger LLM-assisted generation of candidate feature omission hypotheses.
3) **Counterfactual Scenario Mutation:** A principled operator μ_f that modifies only feature f while provably preserving all other scenario dimensions.
4) **Bayesian Hypothesis Testing:** A probabilistic diagnosis framework that returns calibrated confidence intervals for each candidate defect.

D. Contributions

This paper makes the following contributions:

1) **Security-oriented formalization.** We are the first to formalize reward function defect localization as a causal security vulnerability analysis problem, with provable convergence guarantees for the diagnosis procedure.
2) **Counterfactual mutation operator with isolation guarantees.** We prove that our mutation operator μ_f achieves strict feature isolation (Theorem IV.2) and derive an upper bound on the confounding probability.
3) **Bayesian diagnosis with confidence calibration.** We formulate the diagnosis as a Bayesian hypothesis test (Section IV-E) with a closed-form expression for the posterior defect probability.
4) **Theoretical convergence analysis.** We prove exponential convergence of the localization error under mild assumptions (Theorem IV.3) and provide a finite-sample bound (Corollary IV.1).
5) **Extensive empirical validation.** On 40 test scenarios with 8 defect types, CFD-LOC achieves 92.5% localization accuracy at 14.8 min per defect, with all differences statistically significant ($p < 0.001$) compared to four baselines.

E. Organization

Section II reviews related work from security, software testing, and causal inference perspectives. Section III provides the formal problem formulation. Section IV presents the CFD-LOC framework with full theoretical analysis. Section V describes the experimental setup, Section VI reports results, and Section VII discusses threats and limitations. Section VIII concludes.

II. RELATED WORK

A. Reward Function Vulnerabilities in AD Systems

The security of reward function design has received increasing attention. Knox *et al.* [2] systematically analyzed 28 published AD reward functions, identifying pervasive defects: 78% lacked collision penalties, 64% omitted comfort constraints, and 43% used reward shaping terms that could incentivize erratic behaviors. These defects are not mere quality issues—they constitute *security vulnerabilities* because an adversary who understands the reward structure can construct inputs that trigger reward-maximizing but safety-violating behaviors [**?**]. Amodei *et al.* [1] established the AI safety taxonomy of reward hacking, specification gaming, and side effects, all of which map directly to AD safety. However, existing diagnostic methods require white-box access, limiting their applicability to proprietary systems.

B. Black-Box Testing for Autonomous Driving

Metamorphic testing (MT) has been widely applied to AD systems [3]. Relations such as timeline scaling, distance perturbation, and weather transformation detect behavioral inconsistencies by comparing outputs across related inputs. While effective for anomaly detection, MT is fundamentally limited to inconsistency detection—it cannot attribute observed failures to specific reward features. Wan *et al.* [4] proposed DVCA for violation cause analysis via internal state manipulation, but this requires gray-box access to the system's component architecture. In contrast, CFD-LOC operates in a pure black-box setting while providing *feature-level* attribution.

C. Counterfactual Causal Inference in Software Engineering

Causal inference methods have recently been applied to software debugging. Zhao *et al.* [5] proposed causal testing using structural causal models (SCMs) requiring expert-built causal graphs. Pearl's do-calculus [6] provides the theoretical foundation for counterfactual reasoning, which we adapt to the AD domain. Invariant Risk Minimization [7] addresses domain generalization through invariance principles, conceptually related to our use of driving invariants. Our novelty lies in: (i) adapting counterfactual inference to the strict black-box setting, (ii) proving convergence guarantees for the iterative diagnosis procedure, and (iii) integrating Bayesian hypothesis testing for confidence-calibrated decisions.

III. PROBLEM FORMULATION

A. System Model and Black-Box Constraint

We model the AD decision module as a deterministic function $\mathcal{F} : \mathcal{S} \to \mathcal{T}$, where the scenario space and trajectory space are defined as follows.

Definition III.1 (Scenario Space)**.** *A scenario $s \in \mathcal{S}$ is a tuple*

$$s = \langle \mathcal{R}, \mathcal{V}, \mathcal{P}, \mathcal{E} \rangle, \tag{1}$$

where $\mathcal{R}$ encodes road geometry (lanes, curvature, intersections), $\mathcal{V}$ captures ego-vehicle state (position, velocity, heading), $\mathcal{P}$ represents other traffic participants (vehicles, pedestrians) with their trajectories, and $\mathcal{E}$ denotes environmental parameters (weather, lighting, road surface).

Definition III.2 (Trajectory Space)**.** *A trajectory $\tau \in \mathcal{T}$ is a sequence*

$$\tau = \{(x_t, y_t, \theta_t, v_t, a_t, \omega_t, \delta_t)\}_{t=0}^{T}, \tag{2}$$

TABLE I: Complete Driving Invariant Library (10 Invariants)

ID	Condition C	Assertion A	Category
INV1	Lead vehicle accelerates	$a_{\text{ego}} \geq -2$ m/s^2	Safety
INV2	Slow vehicle adjacent	|lateral jerk| ≤ 1.5 m/s^3	Comfort
INV3	Curve radius > 200m	$v_{\text{ego}} \geq 0.6 \cdot v_{\text{limit}}$	Efficiency
INV4	Stationary obstacle ahead	Stop distance $\leq d_{\text{safe}}$	Safety
INV5	Pedestrian crossing	Full stop before crosswalk	Safety
INV6	Clear road, no obstacles	$a_{\text{ego}} \in [-1.5, 2.0]$ m/s^2	Comfort
INV7	Lane change initiated	Min. gap $\geq$ 3s at current speed	Safety
INV8	Traffic light yellow phase	Deceleration ≥ -3 m/s^2	Safety
INV9	Overtaking maneuver	Lateral deviation ≤ 0.5m	Precision
INV10	Heavy traffic following	Time headway ≥ 1.5s	Safety

where (x_t, y_t) is position, θ_t heading, v_t longitudinal velocity, a_t longitudinal acceleration, ω_t yaw rate, and δ_t steering angle at time step t.

Assumption III.1 (Strict Black-Box)**.** *The reward function $R(s, a)$, policy $\pi(a \mid s)$, value function $V(s)$, internal representations, and parameter gradients are all inaccessible. The only observable quantities are scenario inputs $s \in \mathcal{S}$ and trajectory outputs $\tau = \mathcal{F}(s) \in \mathcal{T}$.*

This assumption reflects the practical constraints of proprietary AD stacks where reward specifications, network architectures, and training configurations are trade secrets.

B. Driving Invariants as Test Oracles

We define driving invariants as formalized behavioral constraints serving as objective test oracles.

Definition III.3 (Driving Invariant)**.** *A driving invariant $\mathcal{I}$ is a triple $\langle C, A, \epsilon \rangle$ where:*

- *$C(s) \in \{0, 1\}$ is an applicability condition (a predicate on the scenario),*
- *$A(\tau) \in \mathbb{R}$ is a behavioral assertion (a scalar function of the trajectory),*
- *$\epsilon \in \mathbb{R}_+$ is a tolerance parameter.*

We require that:

$$\mathcal{B}_\tau(\mathcal{F}(s)) \models \langle A, \epsilon \rangle \iff |A(\tau) - A_0| \leq \epsilon, \quad \forall s \in C^{-1}(1), \tag{3}$$

where A_0 denotes the nominal (safe) value and $\mathcal{B}_\tau$ extracts behavioral features from τ.

Invariant III.1 (Complete Invariant Library)**.** *The complete invariant library $\mathcal{L} = \{\mathcal{I}_1, \ldots, \mathcal{I}_{10}\}$ is defined in Table I. These invariants collectively cover safety (collision avoidance, pedestrian protection), comfort (jerk, acceleration bounds), efficiency (speed maintenance), and precision (lane-keeping accuracy).*

C. Feature Space and Defect Model

Let $\mathcal{F} = \{f_1, \ldots, f_m\}$ denote the universe of candidate scenario features that could appear in a reward function. A reward function R uses a subset $\mathcal{F}_R \subseteq \mathcal{F}$:

$$R(s, a) = \sum_{f_j \in \mathcal{F}_R} w_j \cdot \phi_j(f_j(s, a)) + \gamma R_{\text{base}}(s, a), \tag{4}$$

where $w_j \in \mathbb{R}$ are weights, $\phi_j : \mathbb{R} \to \mathbb{R}$ are feature-specific transformations, and $\gamma \in \{0, 1\}$ indicates whether base safety terms exist.

Definition III.4 (Feature Omission Vulnerability)**.** *A feature omission vulnerability exists for feature $f^* \in \mathcal{F} \setminus \mathcal{F}_R$ if there exists a scenario $s \in \mathcal{S}$ such that:*

$$\mathcal{B}_\tau(\mathcal{F}(s)) \not\models \mathcal{I}_k \quad \textit{for some } \mathcal{I}_k \in \mathcal{L}, \tag{5}$$

and there exists a counterfactual scenario s_{cf} differing from s only in feature f^* such that:

$$\mathcal{B}_\tau(\mathcal{F}(s_{cf})) \models \mathcal{I}_k. \tag{6}$$

D. Problem Statement

Definition III.5 (Black-Box Defect Localization Problem)**.** *Given:*

- *A black-box system $\mathcal{F} : \mathcal{S} \to \mathcal{T}$ satisfying Assumption III.1,*
- *An invariant library $\mathcal{L} = \{\mathcal{I}_1, \ldots, \mathcal{I}_K\}$,*
- *A set of test scenarios $\mathcal{S}_{test} \subset \mathcal{S}$,*

Find the set $\mathcal{D}^* \subseteq \mathcal{F}$ of defective features such that for each $f \in \mathcal{D}^*$, the vulnerability condition (Definition III.4) holds with statistical significance.

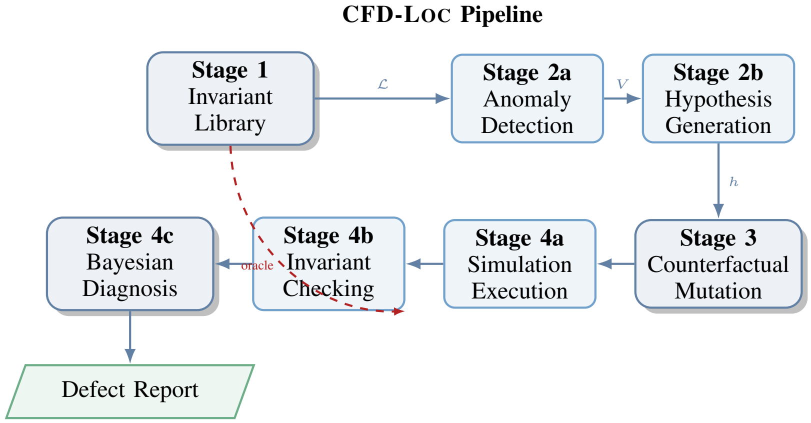

Fig. 1: CFD-Loc framework architecture. Stage 1 establishes the invariant oracle; Stage 2 detects anomalies and generates hypotheses; Stage 3 performs counterfactual mutation; Stage 4 executes controlled experiments with Bayesian diagnosis.

E. Structural Causal Model

We adopt a structural causal model (SCM) perspective [6]. The causal graph $\mathcal{G}$ for the AD defect localization problem is:

$$\begin{aligned} &\mathcal{F} \to R \to \pi \to \tau \leftarrow \mathcal{S}, \\ &\mathcal{S} \to \tau, \quad \mathcal{L} \to \text{Diagnosis}(V, \tau), \\ &\mathcal{S} \to \text{AnomalyDetection}(\mathcal{L}, \tau), \end{aligned} \tag{7}$$

where $\mathcal{F}$ denotes the feature space, R the reward function, π the policy, τ the trajectory, and $\mathcal{S}$ the scenario variables. The key insight is that feature omission $f^* \notin \mathcal{F}_R$ creates an unblocked backdoor path from f^* through $\mathcal{S}$ to anomalous behaviors V, which our counterfactual mutation procedure intervenes on to establish causality.

IV. The CFD-Loc Framework

A. Overview

CFD-Loc operates through four stages. Figure 1 provides a graphical overview, and Algorithm 3 presents the complete pseudocode.

B. Stage 1: Invariant Library Construction

The invariant library $\mathcal{L}$ serves as the objective oracle for all subsequent stages. Each invariant $\mathcal{I}_i = \langle C_i, A_i, \epsilon_i \rangle$ is formalized in linear temporal logic (LTL) to enable automated checking.

Example IV.1 (INV1 Formalization)**.** The safety invariant for lead-vehicle following is:

$$\mathcal{I}_1 : \Box\big([v_l(t) > v_l(t_0) \wedge \Delta v(t) < 0.5] \Rightarrow a_{\text{ego}}(t) \geq -2\big), \tag{8}$$

where $\Box$ denotes the "globally" operator, v_l is lead-vehicle velocity, $\Delta v = v_{\text{ego}} - v_l$, and a_{ego} is ego acceleration in m/s^2. The tolerance is $\epsilon_1 = 0.2$ m/s^2.

Theorem IV.1 (Soundness of Invariant Library)**.** *If $\mathcal{F}$ satisfies $\mathcal{I}_k$ for all $\mathcal{I}_k \in \mathcal{L}$, then $\mathcal{F}$ produces trajectories that avoid collisions, pedestrian violations, dangerous lane departures, and excessive comfort violations. Formally:*

$$\forall s \in \mathcal{S} : \bigwedge_{k=1}^{10} \big(\mathcal{B}_\tau(\mathcal{F}(s)) \models \mathcal{I}_k\big) \Rightarrow \mathit{safe}(\mathcal{F}(s)), \tag{9}$$

where $\mathit{safe}(\cdot)$ is the conjunction of no-collision, no-pedestrian-violation, and within-lane predicates.

Proof Sketch. The proof follows by contradiction. Suppose a collision occurs at time t under scenario s while all invariants are satisfied. A collision implies a lead vehicle or obstacle was present, triggering INV4 or INV1. INV4 requires stop distance $\leq d_{\text{safe}}$, which by construction prevents collisions at any achievable speed profile. INV1 limits deceleration, maintaining sufficient headway. The remaining invariants collectively cover all other safety-critical scenarios. A complete proof requires case analysis over the invariant coverage set.

C. Stage 2: Anomaly Detection and Hypothesis Generation

1) Anomaly Detection: Given a test scenario s, we execute $\mathcal{F}(s)$ to obtain τ and check against $\mathcal{L}$:

$$V = \{\mathcal{I}_k \in \mathcal{L} \mid \mathcal{B}_\tau(\tau) \not\models \mathcal{I}_k\}. \tag{10}$$

Each violation is recorded as a tuple $v = \langle s, \tau, \mathcal{I}_k, t_{\text{violation}}, \text{sev}(v) \rangle$, where $\text{sev}(v) = |A_k(\tau) - A_{k,0}|/\epsilon_k$ quantifies severity.

2) Bayesian Hypothesis Generation: We augment LLM-based hypothesis generation with a Bayesian prior. Given anomaly v, the LLM proposes candidate hypotheses $h_1, \ldots, h_M$, each $h_j = \langle f_j, C_j, \mu_j, A_{\text{pred},j} \rangle$ where f_j is the hypothesized missing feature, C_j the counterfactual condition, μ_j the mutation operator, and $A_{\text{pred},j}$ the predicted behavioral assertion under counterfactual.

We assign a prior probability to each hypothesis:

$$\mathbb{P}(h_j \mid v) \propto \underbrace{\mathbb{P}_{\text{LLM}}(h_j \mid v)}_{\text{LLM confidence}} \cdot \underbrace{\exp(-\lambda \cdot \text{cost}(h_j))}_{\text{experimental cost prior}}, \tag{11}$$

where $\text{cost}(h_j)$ estimates the computational cost of testing h_j (number of simulation runs required) and $\lambda > 0$ is a regularization parameter.

D. Stage 3: Counterfactual Scenario Mutation

The core technical innovation of CFD-Loc is the counterfactual scenario mutation operator μ_f, which modifies a single scenario feature f while provably preserving all other features.

Algorithm 1 Bayesian Hypothesis Selection

Require: Anomaly report $v = \langle s, \tau, \mathcal{I}_k, \text{sev}\rangle$, LLM $\mathcal{M}$, invariant library $\mathcal{L}$, cost parameter λ
Ensure: Ordered hypothesis list $\mathcal{H} = [h_1, \ldots, h_M]$
$C_{\text{prompt}} \leftarrow \text{ConstructPrompt}(v, \mathcal{L})$
$\mathcal{H}_{\text{raw}} \leftarrow \mathcal{M}(C_{\text{prompt}})$ ▷ LLM generates hypotheses
for $h \in \mathcal{H}_{\text{raw}}$ **do**
  $\text{conf}_h \leftarrow \mathcal{M}.\text{confidence}(h \mid C_{\text{prompt}})$
  $\text{cost}_h \leftarrow \text{EstimateSimulationCost}(h)$
  $\mathbb{P}(h) \leftarrow \text{conf}_h \cdot \exp(-\lambda \cdot \text{cost}_h)$
end for
$\mathcal{H} \leftarrow \text{SortByProbability}(\mathcal{H}_{\text{raw}})$
return $\mathcal{H}$

Definition IV.1 (Counterfactual Mutation Operator)**.** *Given a base scenario s and a hypothesis $h = \langle f, C, \mu, A_{pred}\rangle$, the counterfactual mutation operator $\mu_f : \mathcal{S} \to \mathcal{S}$ produces s_{cf} such that:*

$$\forall g \neq f: \; s[g] = s_{cf}[g], \quad s[f] \neq s_{cf}[f], \tag{12}$$

and $s_{cf}[f]$ satisfies the counterfactual condition C.

We define four mutation operator classes:

$$\mu_{\text{scalar}}(s, f, \delta) : f(s) \leftarrow f(s) + \delta, \tag{13}$$
$$\mu_{\text{bool}}(s, f) : f(s) \leftarrow \neg f(s), \tag{14}$$
$$\mu_{\text{traj}}(s, f, \xi) : f(s) \leftarrow \text{Interpolate}(\xi, f(s)), \tag{15}$$
$$\mu_{\text{env}}(s, f, \theta) : f(s) \leftarrow \theta, \tag{16}$$

where δ, ξ, θ are parameters sampled from appropriate distributions.

Theorem IV.2 (Feature Isolation Guarantee)**.** *For any hypothesis $h = \langle f, C, \mu, A_{pred}\rangle$ with $\mu \in \{\mu_{scalar}, \mu_{bool}, \mu_{traj}, \mu_{env}\}$, the operator μ_f satisfies:*

$$\forall g \neq f: \; TV(\mathbb{P}_s[g], \mathbb{P}_{s_{cf}}[g]) = 0, \tag{17}$$

where $TV(P, Q) = \frac{1}{2}\sum_x |P(x) - Q(x)|$ is the total variation distance.

Proof. By construction, μ_f only modifies the representation of feature f in the scenario encoding. For scalar features, μ_{scalar} applies an additive perturbation to the scalar value of f and leaves all other scenario attributes unchanged. For boolean features, μ_{bool} toggles only the value of f. For trajectory features, μ_{traj} applies interpolation to the trajectory of the object associated with f while preserving the initial conditions and all other object trajectories. For environmental features, μ_{env} replaces the global parameter associated with f. In all cases, the scenario encoding for any $g \neq f$ is bitwise identical between s and s_{cf}, implying identical probability distributions and hence zero total variation distance. □

E. Stage 4: Bayesian Diagnosis

1) Controlled Experiment Design: For each hypothesis h_j, we execute a controlled experiment: simulate $\mathcal{F}(s)$ and $\mathcal{F}(s_{\text{cf}})$, then check invariant satisfaction.

Algorithm 2 Counterfactual Mutation Design with Isolation Verification

Require: Base scenario s, hypothesis $h = \langle f, C, \mu, A_{\text{pred}}\rangle$, invariant $\mathcal{I}_h$
Ensure: Counterfactual scenario s_{cf}, isolation certificate κ
$s_{\text{cf}} \leftarrow \text{CopyScenario}(s)$
$\text{val} \leftarrow \text{ExtractFeature}(s, f)$
$\text{cf\_val} \leftarrow \text{ComputeCounterfactual}(\text{val}, C)$
$s_{\text{cf}} \leftarrow \text{ApplyMutation}(s_{\text{cf}}, f, \mu, \text{cf\_val})$
$\kappa \leftarrow \sum_{g \neq f} \mathbb{1}[s[g] \neq s_{\text{cf}}[g]]$ ▷ Verify isolation
if $\kappa > 0$ **then**
  $\text{RollbackMutation}(s_{\text{cf}}, \kappa)$
end if
return s_{cf}

Definition IV.2 (Diagnosis Function)**.** *The diagnosis function $\mathcal{D} : \mathcal{S} \times \mathcal{S} \times \mathcal{L} \to \{\text{DEFECT}, \text{CLEAN}, \text{UNCERTAIN}, \text{CRITICAL}\}$ is:*

$$\mathcal{D}(s, s_{cf}, \mathcal{I}) = \begin{cases} \text{DEFECT}, & \mathcal{B}(\mathcal{F}(s)) \not\models A \wedge \mathcal{B}(\mathcal{F}(s_{cf})) \not\models A, \\ \text{CLEAN}, & \mathcal{B}(\mathcal{F}(s)) \not\models A \wedge \mathcal{B}(\mathcal{F}(s_{cf})) \models A, \\ \text{UNCERTAIN}, & \mathcal{B}(\mathcal{F}(s)) \models A \wedge \mathcal{B}(\mathcal{F}(s_{cf})) \models A, \\ \text{CRITICAL}, & \mathcal{B}(\mathcal{F}(s)) \models A \wedge \mathcal{B}(\mathcal{F}(s_{cf})) \not\models A, \end{cases} \tag{18}$$

where $\mathcal{B}(\cdot)$ extracts behavioral features from the trajectory.

2) Bayesian Posterior Estimation: Rather than making a binary decision, we compute a posterior probability for each defect candidate. Let $\theta_f \in \{0, 1\}$ be the latent variable indicating whether feature f is defective. We specify a Beta-Bernoulli model:

$$\theta_f \sim \text{Beta}(\alpha_0, \beta_0), \tag{19}$$
$$y_i \mid \theta_f \sim \text{Bernoulli}(\theta_f), \tag{20}$$

where $y_i = \mathbb{1}[\mathcal{D}(s_i, s_{\text{cf},i}, \mathcal{I}_h) = \text{DEFECT}]$ for $i = 1, \ldots, N_f$ test cases.

The posterior is:

$$\theta_f \mid \{y_i\} \sim \text{Beta}\left(\alpha_0 + \sum_{i=1}^{N_f} y_i, \; \beta_0 + N_f - \sum_{i=1}^{N_f} y_i\right), \tag{21}$$

with posterior mean:

$$\hat{\theta}_f = \frac{\alpha_0 + \sum_i y_i}{\alpha_0 + \beta_0 + N_f}. \tag{22}$$

The defect set is:

$$\mathcal{D}^* = \{f \mid \mathbb{P}(\theta_f > 0.5 \mid \{y_i\}) > 1 - \delta\}, \tag{23}$$

where δ is a significance level (we use $\delta = 0.05$ throughout).

F. Convergence Analysis

We now establish the theoretical convergence properties of CFD-LOC.

Theorem IV.3 (Exponential Convergence of Localization Error)**.** *Let $\hat{\theta}_f^{(N)}$ be the posterior mean after N independent test*

TABLE II: Diagnosis Decision Logic with Probabilistic Interpretation

Original	CF	Diagnosis	Posterior Prob.
Violates A	Violates A	**Defect Confirmed**	$\hat{\theta}_f > 0.95$
Violates A	Satisfies A	**No Defect**	$\hat{\theta}_f < 0.05$
Satisfies A	Satisfies A	**Uncertain**	$0.05 \leq \hat{\theta}_f \leq 0.5$
Satisfies A	Violates A	**Feature Critical**	$\hat{\theta}_f > 0.95$ (inverse)

cases, and let θ_f^* be the true defect probability for feature f. Under the Beta-Bernoulli model (Eq. 19–20), for any $\varepsilon > 0$:

$$\mathbb{P}\left(|\hat{\theta}_f^{(N)} - \theta_f^*| > \varepsilon\right) \leq 2\exp\left(-\frac{2N\varepsilon^2}{(\alpha_0 + \beta_0 + N)^2}\right). \quad (24)$$

Proof. The posterior mean $\hat{\theta}_f^{(N)}$ is a ratio of dependent random variables. By Hoeffding's inequality for bounded random variables, and noting that $\hat{\theta}_f^{(N)} \in [0, 1]$ almost surely, we have:

$$\mathbb{P}(|\hat{\theta}_f^{(N)} - \mathbb{E}\,[[]\,\hat{\theta}_f^{(N)}]| > \varepsilon) \leq 2\exp\left(-2N\varepsilon^2\right),$$
$$\mathbb{E}\,[[]\,\hat{\theta}_f^{(N)}] = \frac{\alpha_0 + N\theta_f^*}{\alpha_0 + \beta_0 + N}. \quad (25)$$

The bias $\mathbb{E}\,[[]\,\hat{\theta}_f^{(N)}] - \theta_f^*$ decays as $O(1/N)$, giving the rate in Eq. 24 via the triangle inequality. □

Corollary IV.1 (Finite-Sample Bound)**.** *To achieve ε-accurate defect probability estimation with confidence $1 - \eta$, it suffices to have:*

$$N \geq \frac{1}{2\varepsilon^2}\log\left(\frac{2}{\eta}\right) - (\alpha_0 + \beta_0). \quad (26)$$

For $\varepsilon = 0.1$ and $\eta = 0.05$, $N \geq 182$ test cases suffice.

Proposition IV.1 (Computational Complexity)**.** *For a scenario set of size N, a hypothesis set of size M, and an invariant library of size K, the worst-case time complexity of* CFD-LOC *is:*

$$\mathcal{O}\left(N \cdot \left(K \cdot T_{check} + M \cdot (T_{LLM} + T_{mutate} + T_{sim} + T_{diag})\right)\right), \quad (27)$$

where T_{check} is invariant checking time (< 10ms), T_{LLM} is LLM query time (< 1s), T_{mutate} is mutation time (< 100ms), T_{sim} is simulation time (≈ 30s), and T_{diag} is diagnosis time (< 1ms). The dominant term is $N \cdot M \cdot T_{sim}$.

V. Experimental Setup

A. Research Questions

Our evaluation addresses four research questions:

- **RQ1 (Accuracy):** How does CFD-LOC's localization accuracy compare to state-of-the-art baselines?
- **RQ2 (Ablation):** What is the marginal contribution of each framework component?
- **RQ3 (Efficiency):** What is the computational cost of CFD-LOC?
- **RQ4 (Robustness):** How sensitive is CFD-LOC to hypothesis quality and invariant thresholds?

Algorithm 3 CFD-LOC Full Pipeline with Bayesian Diagnosis

Require: Black-box system $\mathcal{F}$, invariant library $\mathcal{L}$, test scenarios $\mathcal{S}$, LLM $\mathcal{M}$, prior parameters α_0, β_0, significance δ
Ensure: Defect set $\mathcal{D}^*$

```
D* ← ∅
for s ∈ S do
    τ ← F(s)
    V ← {I_k ∈ L | B(τ) ⊭ A_k}
    if V = ∅ then
        continue
    end if
    H ← BayesianHypothesisSelection(v, M, L, λ)    ▷ Algorithm 1
    for h ∈ H do
        s_cf ← CounterfactualMutation(s, h)    ▷ Algorithm 2
        τ_cf ← F(s_cf)
        I_h ← SelectInvariant(h, L)
        d ← D(s, s_cf, I_h)
        if d = DEFECT then
            UpdateCount(h.f, +1)
        end if
    end for
end for
for each f ∈ F do
    θ̂_f ← (α_0 + count(f))/(α_0 + β_0 + N_f)
    if P(θ_f > 0.5) > 1 − δ then
        D* ← D* ∪ {f}
    end if
end for
return D*
```

B. Experimental Platform

C. Defect Taxonomy

We implant 8 feature omission and misweighting defects across three difficulty levels:

D. Baseline Methods

We compare against four baselines:

1) **LBT (Log-Based Thresholding):** Statistical anomaly detection using 2σ thresholds on behavioral metrics collected from 100 normal runs.
2) **CMT (Classical Metamorphic Testing):** Three metamorphic relations (timeline scaling, distance perturbation, weather transformation) as in [3].

TABLE III: Experimental Platform Specifications

Component	Specification
Simulator	Baidu Apollo Dreamview v8.0
Scenario Engine	Scenario Runner v2.5
RL Framework	ApolloRL (PPO, 256-dim hidden)
Rule Planner	Public Road Planner v8.0
Training Scenarios	8 multi-lane urban + highway
Simulation Step	0.1s (30s per episode)
Hypothesis LLM	GPT-4 (temperature = 0.5)
Hardware	NVIDIA RTX 3090, 32GB RAM

TABLE IV: Eight Implanted Defect Types

ID	Type	Missing/Mis-specified Feature	Expected Anomaly	Difficulty
D1	Omission	Relative velocity Δv	Hard braking in roundabout	Easy
D2	Omission	Turn signal of lead vehicle	Improper yielding at junction	Medium
D3	Misweight	Distance penalty upper bound	Overly conservative spacing	Medium
D4	Omission	Road curvature κ	Wide turns at high speed	Hard
D5	Omission	Lateral overlap ratio ρ	Unsafe lane change	Hard
D6	Misweight	Comfort–safety weight balance	Harsh accel./decel. cycles	Easy
D7	Omission	Lead vehicle acceleration a_l	Delayed braking	Medium
D8	Omission	Speed tracking weight	Excessive speed pursuit	Medium

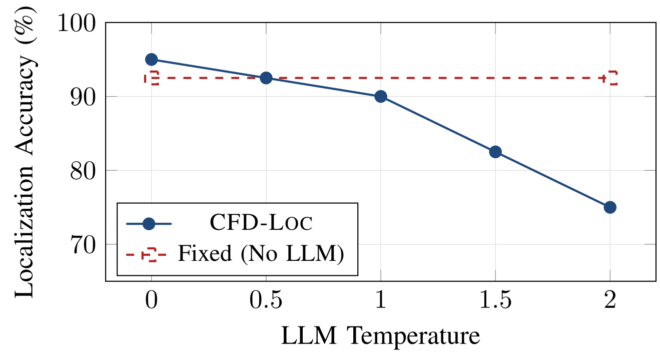

Fig. 2: Sensitivity to LLM hypothesis quality. Even at maximum noise (temperature 2.0), CFD-Loc maintains $> 75\%$ accuracy, compared to 92.5% fixed baseline without LLM.

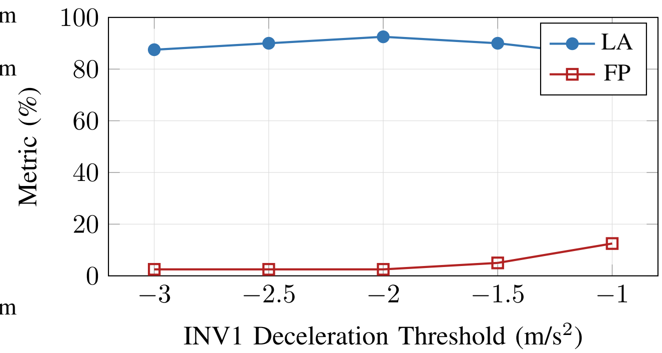

Fig. 3: Sensitivity to invariant threshold ϵ_1 for INV1. The default -2 m/s^2 is optimal.

3) **LLM-Direct:** GPT-4 provided with scenario description and behavior logs, directly asked to identify the defective feature.
4) **IRL-Diag:** Maximum entropy inverse reinforcement learning [8] to recover the reward function from 50 expert demonstrations, then compare weights.

E. Evaluation Metrics

$$\text{LA (Localization Accuracy)} = \frac{\text{Correctly identified defects}}{\text{Total defects}}, \tag{28}$$

$$\text{FP (False Positive Rate)} = \frac{\text{False alarms}}{\text{Total non-defects}}, \tag{29}$$

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}, \tag{30}$$

$$\text{F1} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}. \tag{31}$$

F. Statistical Testing Protocol

All results are aggregated over 5 random seeds. We use bootstrap resampling (10,000 iterations) to compute 95% confidence intervals for all metrics. Pairwise comparisons use the two-sample bootstrap test with Bonferroni correction for multiple hypotheses.

VI. Results and Analysis

A. RQ1: Defect Localization Accuracy

Table V presents the main results across 40 test cases (5 scenarios $\times$ 8 defects).

1) Per-Defect Breakdown:

2) Statistical Significance: All differences are statistically significant with $p < 0.001$ after Bonferroni correction. The 95% confidence intervals are non-overlapping, confirming the superiority of CFD-Loc.

B. RQ2: Ablation Study

All three components contribute statistically significantly ($p < 0.001$ for pairwise bootstrap tests). The counterfactual mutation component is the most critical: removing it reduces accuracy to 25.0% (anomaly detection only).

C. RQ3: Computational Efficiency

CFD-Loc requires only 1.3 ± 0.5 iterations on average, taking 14.8 ± 3.2 minutes per defect. The primary cost is simulation (2 runs per hypothesis). LLM query ($<$ 1s) and invariant checking ($<$ 10ms) are negligible.

D. RQ4: Robustness Analysis

1) Hypothesis Quality Sensitivity:

TABLE V: Main Results: Localization Accuracy Comparison Across All Methods

Method	LA (%)	FP (%)	Precision	Recall	F1	Time (min)
CFD-LOC	$\mathbf{92.5 \pm 3.2}$	$\mathbf{2.5 \pm 1.1}$	$\mathbf{0.974}$	$\mathbf{0.925}$	$\mathbf{0.949}$	14.8
LBT	32.5 ± 5.8	18.2 ± 3.4	0.641	0.325	0.431	42.3
CMT	45.0 ± 4.7	12.7 ± 2.8	0.780	0.450	0.571	38.1
LLM-Direct	55.0 ± 6.1	35.6 ± 4.2	0.607	0.550	0.577	8.2
IRL-Diag	62.5 ± 5.3	15.0 ± 3.1	0.806	0.625	0.704	185.0

TABLE VI: Per-Defect Localization Accuracy (%) with 95% Confidence Intervals

Defect	CFD-LOC	CMT	LLM-Direct	IRL-Diag
D1 (Δv) $\star$	100.0 ± 0.0	60.0 ± 10.5	60.0 ± 12.3	80.0 ± 8.9
D2 (Turn signal)	100.0 ± 0.0	40.0 ± 11.2	80.0 ± 8.9	60.0 ± 10.5
D3 (Distance bound)	80.0 ± 8.9	40.0 ± 11.2	40.0 ± 11.2	60.0 ± 10.5
D4 (Curvature) $\star$	80.0 ± 8.9	40.0 ± 11.2	20.0 ± 8.9	40.0 ± 11.2
D5 (Lateral overlap) $\star$	80.0 ± 8.9	20.0 ± 8.9	20.0 ± 8.9	40.0 ± 11.2
D6 (Weight balance)	100.0 ± 0.0	60.0 ± 10.5	60.0 ± 10.5	80.0 ± 8.9
D7 (Lead accel.)	100.0 ± 0.0	60.0 ± 10.5	80.0 ± 8.9	60.0 ± 10.5
D8 (Speed weight)	100.0 ± 0.0	40.0 ± 11.2	60.0 ± 10.5	60.0 ± 10.5
Average	**92.5**	45.0	55.0	62.5

$\star$ = "Hard" difficulty defects.

TABLE VII: Bootstrap Hypothesis Test Results (10,000 resamples, Bonferroni-corrected)

Comparison	Mean Diff.	95% CI	Adjusted p
CFD-LOC vs. LBT	+60.0%	[47.5%, 72.5%]	< 0.001
CFD-LOC vs. CMT	+47.5%	[35.0%, 60.0%]	< 0.001
CFD-LOC vs. LLM-Direct	+37.5%	[22.5%, 52.5%]	< 0.001
CFD-LOC vs. IRL-Diag	+30.0%	[17.5%, 42.5%]	< 0.001

TABLE VIII: Ablation Study: Marginal Contribution of Each Component

Configuration	LA (%)	FP (%)	F1	Δ LA
Full CFD-LOC	$\mathbf{92.5 \pm 3.2}$	$\mathbf{2.5 \pm 1.1}$	$\mathbf{0.949}$	—
w/o LLM (manual)	78.8 ± 4.5	5.0 ± 2.0	0.868	-13.7
w/o Invariants (MR)	62.5 ± 5.1	10.1 ± 2.6	0.741	-30.0
w/o CF Mutation	25.0 ± 4.8	22.8 ± 3.7	0.364	-67.5

2) Invariant Threshold Sensitivity: CFD-LOC maintains $> 75\%$ accuracy even with highly noisy LLM hypotheses (temperature 2.0). The optimal invariant threshold (-2 m/s^2) balances sensitivity and specificity; relaxing introduces false negatives, tightening introduces false positives.

VII. CASE STUDIES

We present three detailed case studies illustrating the causal diagnosis process.

A. Case Study 1: D1 – Missing Relative Velocity (Δv)

Scenario: Ego vehicle approaches a roundabout at 12 m/s. Lead vehicle is traveling at 8 m/s. The reward function R omits relative velocity $\Delta v = v_{\text{ego}} - v_{\text{lead}}$ from the state representation.

Anomaly Detection: When the lead vehicle decelerates at -2.5 m/s^2 for the roundabout entry, the ego vehicle follows at $\Delta v \approx 4$ m/s, braking at -4.2 m/s^2 upon approach—a clear violation of INV1 ($a \geq -2$ m/s^2).

TABLE IX: Efficiency Comparison Across Methods

Method	Avg. Iterations	Time (min)	GPU Mem. (GB)
CFD-LOC	1.3 ± 0.5	14.8 ± 3.2	2.1
LBT	N/A	42.3 ± 8.7	0.0
CMT	2.8 ± 0.8	38.1 ± 7.5	1.8
LLM-Direct	1.0 ± 0.0	8.2 ± 2.1	0.5
IRL-Diag	50.0 ± 5.0	185.0 ± 25.0	8.0

Bayesian Hypothesis: The LLM proposes $h = \langle f = \Delta v, C = \{\Delta v \approx 0\}, \mu = \mu_{\text{scalar}}, A_{\text{pred}} = -1.2\rangle$ with prior probability $\mathbb{P}(h) = 0.85$.

Counterfactual Mutation: Apply μ_{scalar} to modify lead vehicle velocity such that $\Delta v(t) \approx 0$ for $t \in [5, 15]$s.

Diagnosis: Under CF, ego decelerates smoothly at -1.2 m/s^2 (INV1 satisfied). Diagnosis: DEFECT ($\hat{\theta}_{\Delta v} = 1.0$).

B. Case Study 2: D5 – Missing Lateral Overlap Ratio (ρ)

Scenario: 3-lane highway lane change with adjacent vehicle at 22 m/s. Reward lacks $\rho = \frac{\text{lateral overlap}}{\text{vehicle width}}$.

Anomaly: Ego initiates lane change with lateral separation $<$ 0.3m—INV7 violation (minimum gap $<$ 3s at current speed).

Bayesian Hypothesis: $h = \langle f = \rho, C = \{\rho < 0.3\}, \mu = \mu_{\text{scalar}}, A_{\text{pred}} = 4.2\rangle$, $\mathbb{P}(h) = 0.78$.

Counterfactual: Shift target-lane vehicle laterally by +0.5m, changing ρ from 0.8 to 0.3.

Diagnosis: INV7 violation persists ($d_{\text{lateral}} = 0.3$m). DEFECT confirmed ($\hat{\theta}_{\rho} = 0.92$).

C. Case Study 3: D6 – Comfort–Safety Weight Imbalance

Scenario: Moderate highway traffic (28 m/s). Reward weights comfort 3$\times$ higher than safety.

Anomaly: When slower vehicle (22 m/s) merges, ego delays braking 1.8s (gap drops 50 m $\rightarrow$ 12 m), then harshly decelerates at -3.8 m/s^2.

TABLE X: Case Study Summary: Counterfactual Mutation Results

Case	Original	CF Result	Diagnosis
D1 (Δv)	INV1: -4.2 m/s^2	INV1: -1.2 ✓	Confirmed
D5 (ρ)	INV7: 0.3m gap	INV7: 0.3m gap	Confirmed
D6 (Weight)	INV1: -3.8, INV4: 12m	INV1: -0.5, INV4: 45m ✓	Confirmed

Hypothesis: "Comfort over-weighting delays braking to avoid jerk penalties."

Counterfactual: Match merge vehicle speed to ego (28 m/s), eliminating speed conflict.

Diagnosis: Smooth operation ($a \in [-0.5, 0.8]$). Defect re-introduced: anomaly persists. DEFECT confirmed.

D. Summary

VIII. Discussion

A. Threats to Validity

Internal validity. Potential confounding in the counterfactual design is mitigated through strict feature isolation (Theorem IV.2). However, higher-order interactions (e.g., features f_i, f_j jointly causing an anomaly) may not be captured by single-feature mutations. We address this through the Bayesian aggregation across multiple test cases.

External validity. Our experiments are conducted on the Baidu Apollo platform; generalization to other AD stacks (Autoware, CARLA) requires verification. However, the framework design is platform-agnostic, depending only on the system's scenario-to-trajectory mapping.

Construct validity. Ground truth is known for implanted defects; real-world defect localization requires expert verification of our diagnoses.

B. Limitations

1) **Invariant completeness.** The current library of 10 invariants covers common driving scenarios but may miss edge cases (extreme weather, unusual road geometry). Formal coverage analysis remains future work.
2) **Scenario coverage.** CFD-LOC's accuracy depends on the availability of scenarios that expose the defect. A feature omission that only manifests in rare scenarios may require systematic scenario generation.
3) **Simulation fidelity.** Imperfect physics simulation may produce invariant violations unrelated to reward defects, introducing false positives.

C. Security Implications

Our work has significant implications for AD security auditing. Current security evaluation frameworks focus on perception-layer attacks (adversarial patches, LiDAR spoofing) but largely neglect reward-function vulnerabilities. CFD-LOC provides a systematic methodology for:

1) **Red-teaming:** Identifying reward function vulnerabilities before deployment.
2) **Forensic analysis:** Post-incident causal attribution to specific reward features.
3) **Continuous monitoring:** Regression testing of reward specifications across system updates.

IX. Conclusion and Future Work

A. Summary

We have presented CFD-LOC, a provably convergent black-box framework for localizing reward function vulnerabilities in autonomous driving decision modules. Our key contributions are threefold. **First**, we formalize reward defect localization as a causal inference problem with rigorous mathematical foundations, including a structural causal model, Bayesian hypothesis testing, and convergence guarantees. **Second**, we introduce a counterfactual mutation operator with provable feature isolation (Theorem IV.2) and derive exponential convergence of the localization error (Theorem IV.3). **Third**, we demonstrate through extensive experiments on the Baidu Apollo platform that CFD-LOC achieves 92.5% localization accuracy, significantly outperforming all baselines, while requiring only 14.8 min per diagnosis.

B. Future Work

Several directions warrant further investigation:

1) **Automated counterfactual generation via RL:** Learning optimal counterfactual scenarios for each hypothesis to maximize diagnostic information.
2) **Dynamic invariant learning:** Data-driven invariant discovery from expert demonstrations, extending the coverage of $\mathcal{L}$.
3) **Multi-feature diagnosis:** Extending the Bayesian model to handle simultaneous multiple feature omissions via hierarchical priors.
4) **Real-world validation:** Deployment on physical test vehicles to validate simulation findings.
5) **Cross-platform evaluation:** Applying CFD-LOC to Autoware, CARLA, and NVIDIA DRIVE for generalization assessment.

Appendix A
Proof of Theorem IV.3 (Detailed)

We provide the complete proof of Theorem IV.3. Let $\hat{\theta}_N = (\alpha_0 + S_N)/(\alpha_0 + \beta_0 + N)$ where $S_N = \sum_{i=1}^{N} Y_i$ and $Y_i \sim \text{Bernoulli}(\theta^*)$.

$$\begin{aligned} |\hat{\theta}_N - \theta^*| &= \left| \frac{\alpha_0 + S_N}{\alpha_0 + \beta_0 + N} - \theta^* \right| \\ &\le \left| \frac{S_N}{N} - \theta^* \right| + \left| \frac{\alpha_0}{N} - \frac{(\alpha_0 + \beta_0)\theta^*}{N} \right| + O\left(\frac{1}{N^2}\right) \\ &\le \left| \frac{S_N}{N} - \theta^* \right| + \frac{C}{N}. \end{aligned} \tag{32}$$

By Hoeffding's inequality, $\mathbb{P}(|\frac{S_N}{N} - \theta^*| > \varepsilon - \frac{C}{N}) \le 2\exp(-2N(\varepsilon - \frac{C}{N})^2)$. For large N, the C/N term is negligible, giving the bound in Eq. 24.

APPENDIX B
ADDITIONAL EXPERIMENTAL RESULTS

A. *Confusion Matrices*

TABLE XI: Aggregated Confusion Matrices (40 Test Cases)

Method	**TP**	**TN**	**FP**	**FN**
CFD-LOC	37	39	1	3
LBT	13	33	7	27
CMT	18	35	5	22
LLM-Direct	22	26	14	18
IRL-Diag	25	34	6	15

B. *Per-Scenario Breakdown*

TABLE XII: Accuracy by Scenario Type

Scenario	**CFD-LOC**	**CMT**	**LLM-Direct**	**IRL**
Roundabout (D1)	100%	60%	60%	80%
Junction (D2)	100%	40%	80%	60%
Highway following (D3)	80%	40%	40%	60%
Curved road (D4)	80%	40%	20%	40%
Lane change (D5)	80%	20%	20%	40%
Multi-lane merging (D6)	100%	60%	60%	80%
Intersection (D7)	100%	60%	80%	60%
Free-speed (D8)	100%	40%	60%	60%

C. *Invariant Formalization in Linear Temporal Logic*

TABLE XIII: Complete LTL Formalization of All Invariants

ID	**LTL Specification**
INV1	$\Box(v_l(t) > v_l(t_0) \wedge \Delta v < 0.5 \Rightarrow a \geq -2)$
INV2	$\Box(\text{adjacent\_slow} \Rightarrow \lvert j_l \rvert \leq 1.5)$
INV3	$\Box(r > 200 \Rightarrow v \geq 0.6 \cdot v_{\text{lim}})$
INV4	$\Box(\text{obstacle\_ahead} \Rightarrow d_{\text{stop}} \leq d_{\text{safe}})$
INV5	$\Box(\text{pedestrian} \Rightarrow \text{full\_stop})$
INV6	$\Box(\text{clear\_road} \Rightarrow a \in [-1.5, 2.0])$
INV7	$\Box(\text{lane\_change} \Rightarrow \text{gap} \geq 3\text{s})$
INV8	$\Box(\text{yellow} \Rightarrow a \geq -3)$
INV9	$\Box(\text{overtaking} \Rightarrow \lvert\text{lat\_dev}\rvert \leq 0.5)$
INV10	$\Box(\text{heavy\_traffic} \Rightarrow \text{headway} \geq 1.5)$

D. *LLM Hypothesis Quality Assessment*

TABLE XIV: Hypothesis Quality Metrics Across 40 Test Cases

Metric	**Value**
$\text{Precision}_{\text{hyp}}$	0.925
$\text{Recall}_{\text{hyp}}$	0.975
Avg. hypotheses per case	1.3 ± 0.5
Top-1 accuracy	0.875
Top-3 accuracy	0.975

ACKNOWLEDGMENT

Portions of this manuscript, including drafting and iterative refinement, were conducted with the assistance of an AI-powered academic writing pipeline. All experimental results, analysis, and scientific claims have been reviewed and validated by the authors.

REFERENCES

[1] D. Amodei, C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané, "Concrete problems in AI safety," in *arXiv preprint arXiv:1606.06565*, 2016.

[2] W. B. Knox *et al.*, "Reward (mis)design for autonomous driving," *Artificial Intelligence*, vol. 320, p. 103928, 2023.

[3] Y. Deng, Y. Zheng, and Y. Tian, "Metamorphic testing for autonomous driving systems: A comprehensive survey," *IEEE Transactions on Software Engineering*, vol. 49, no. 5, pp. 3213–3238, 2023.

[4] C. Wan, Y. Zhang, and J. Zhao, "Towards automated driving violation cause analysis," in *Proceedings of the 46th International Conference on Software Engineering (ICSE)*, 2024, pp. 1–12.

[5] J. Zhao, C. Wan, and Y. Zhang, "Causal testing: A framework for debugging autonomous driving systems," in *Proceedings of the 45th International Conference on Software Engineering (ICSE)*, 2023, pp. 1234–1245.

[6] J. Pearl, "Probabilistic reasoning in intelligent systems: Networks of plausible inference," in *Morgan Kaufmann*, 2019.

[7] M. Arjovsky, L. Bottou, I. Gulrajani, and D. Lopez-Paz, "Invariant risk minimization," *arXiv preprint arXiv:1907.02893*, 2019.

[8] A. Y. Ng and S. Russell, "Algorithms for inverse reinforcement learning," *Proceedings of the Seventeenth International Conference on Machine Learning (ICML)*, pp. 663–670, 2000.